

Database Querying Through Relational Algebra

ACS 575: Database Systems

Instructor: Dr. Jin Soung Yoo, Professor
Department of Computer Science
Purdue University Fort Wayne

References

- ❖ Elmasri et al., Fundamentals of Database Systems, Ch4

Relational Query Languages

- Query languages: Allow manipulation and retrieval of data from a database.
- Relational model supports simple, powerful QLs:
 - Strong formal foundation based on logic.
 - Allows for much optimization.
- Query Languages **!=** programming languages!
 - QLs support easy, efficient access to large data sets.
 - QLs not expected to be “Turing complete”.
 - QLs not intended to be used for complex calculations.

Relational Query Languages

- Two mathematical Query Languages form the basis for “real” languages (e.g. SQL), and for implementation:
 - **Relational Algebra**
 - A procedure query language that can be used to tell the DBMS how to build a new relation from one or more relations in databases
 - **Relational Calculus**
 - A non-procedural language that can be used to formulate the definition of a relation in terms of one or more relations
 - Provides only the description for the query but does not provide the methods to solve it. E.g., $\{t \mid t \in \text{loan} \wedge t[\text{amount}] \geq 10000\}$
- Relational Algebra and Relational Calculous mainly provides theoretical foundation for relational databases and SQL

Relational Algebra

- A theoretical language with operations that work on one or more relations to define another relation without changing the original relation(s).
- Relational Algebra collects instances of relations as input and gives occurrences of relations as output.
 - Each operation takes one or more relations as its operand(s) and generates another relation as its result, e.g., $R1 \cup R2 = R3$
- Relational algebra query operations are performed recursively on a relation.
 - The output from one operation can become the input to another operation.

Basic relational Algebra Operations

- Unary Relational Operations
 - *Selection* (symbol: σ)
 - *Projection* (symbol: π)
 - *Rename* (symbol: ρ)
- Relational Algebra Operations from Set Theory
 - *Union* (\cup)
 - *Intersection* (\cap)
 - *Difference* ($-$)
 - *Cartesian Product* (\times)
- Binary Relational Operations
 - *Join* (\bowtie)
 - *Division*

Example Relations

□ Relation schema

Sailors (sid, sname, rating, age)

Boats (bid, bname, color)

Reserves (sid, bid, day)

□ Example instances

Sailors (*S*)

<u>sid</u>	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

Reserves (*R*)

<u>sid</u>	<u>bid</u>	<u>day</u>
22	101	10/10/96
58	103	11/12/96

Boats (*B*)

<u>bid</u>	bname	color
101	Interlake	blue
102	Interlake	red
103	Clipper	green

Project (π)

- The PROJECT operation eliminates all attributes of the input relation but those mentioned in the projection list.
- The project method defines a relation that contains a vertical subset of relation
 - Schema of result contains exactly the fields in the projection list, with the same names that they had in the (only) input relation.
- This operators is useful for deleting unwanted columns from relation

- E.g.,

<u>sid</u>	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

sname	rating
yuppy	9
lubber	8
guppy	5
rusty	10

age
35.0
55.5
35.0
35.0

Select (σ)

- The SELECT operation is used for selecting a subset of the tuples according to a given selection condition
 $\sigma_p(r)$, where r stands for relation which is the (alias) name of the table, and p is a propositional logic for row selection.
 - e.g., $\sigma_{rating=8 \text{ and } age>50}(S)$
- The schema of result is identical to schema of input relation
- The result relation can be the *input* for another relational algebra operation

Sailors (S)

<u>sid</u>	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

$\sigma_{rating>8}(S)$

sid	sname	rating	age
28	yuppy	9	35.0
58	rusty	10	35.0

$\pi_{sname, rating}(\sigma_{rating>8}(S))$

sname	rating
yuppy	9
rusty	10

Union, Intersection, Set-Difference

□ UNION(\cup)

$$R_1 \cup R_2 = \{r \mid r \in R_1 \text{ or } r \in R_2\}$$

- It performs binary union between two given relations

□ INTERESECT(\cap)

$$R_1 \cap R_2 = \{r \mid r \in R_1 \text{ and } r \in R_2\}$$

- It finds all the tuples that are present in both R_1 and R_2 .

□ Set Difference ($-$)

$$R_1 - R_2$$

- It finds all the tuples that are present in R_1 but not in R_2 .

Union, Intersection, Set-Difference

- All of these operations take two input relations, which must be *union-compatible*:
 - The same number of attributes
 - Attribute domains need to be compatible
 - Duplicate tuples should be automatically removed

sid	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0
44	guppy	5	35.0
28	yuppy	9	35.0

$S1 \cup S2$

$S1$

sid	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

$S2$

sid	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

$S1 \cap S2$

sid	sname	rating	age
31	lubber	8	55.5
58	rusty	10	35.0

$S1 - S2$

sid	sname	rating	age
22	dustin	7	45.0

Cartesian Product (×)

- Cartesian Product is an operation used to merge columns from two relations information of two relations into one.
 $S \times R = \{r\ s \mid r \in R\ or\ s \in S\}$
 - Each row of S is paired with each row of R .
 - *Result schema* has one attribute per attribute of S and R
 - *Conflict*: Both S and R have a field called *sid*.
- Generally, Cartesian product is never a meaningful operation when it performs alone.

S

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

R

<u>sid</u>	<u>bid</u>	<u>day</u>
22	101	10/10/96
58	103	11/12/96

$S \times R$

(sid)	sname	rating	age	(sid)	bid	day
22	dustin	7	45.0	22	101	10/10/96
22	dustin	7	45.0	58	103	11/12/96
31	lubber	8	55.5	22	101	10/10/96
31	lubber	8	55.5	58	103	11/12/96
58	rusty	10	35.0	22	101	10/10/96
58	rusty	10	35.0	58	103	11/12/96

Cross Product

- Cartesian Product becomes meaningful when it is followed by other operations. E.g., $\sigma_{sid=22} (S \times R)$
- It is called Cross Product or Cross Join

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

S

<u>sid</u>	<u>bid</u>	<u>day</u>
22	101	10/10/96
58	103	11/12/96

R

(sid)	sname	rating	age	(sid)	bid	day
22	dustin	7	45.0	22	101	10/10/96
22	dustin	7	45.0	58	103	11/12/96
31	lubber	8	55.5	22	101	10/10/96
31	lubber	8	55.5	58	103	11/12/96
58	rusty	10	35.0	22	101	10/10/96
58	rusty	10	35.0	58	103	11/12/96

S x *R*

(sid)	sname	rating	age	(sid)	bid	day
22	dustin	7	45.0	58	103	11/12/96
31	lubber	8	55.5	58	103	11/12/96

← $\sigma_{S.sid < R.sid} (S \times R)$

Join (\bowtie)

- Join operation is essentially a Cartesian product followed by a selection criterion. $R \bowtie_c S = \sigma_c(R \times S)$
 - The result schema is same as that of cross-product.
 - Fewer tuples than cross-product, might be able to compute more efficiently

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

S

<u>sid</u>	<u>bid</u>	<u>day</u>
22	101	10/10/96
58	103	11/12/96

R

(sid)	sname	rating	age	(sid)	bid	day
22	dustin	7	45.0	58	103	11/12/96
31	lubber	8	55.5	58	103	11/12/96

$S \bowtie_{S.sid < R.sid} R$

(sid)	sname	rating	age	(sid)	bid	day
22	dustin	7	45.0	22	101	10/10/96
22	dustin	7	45.0	58	103	11/12/96
31	lubber	8	55.5	22	101	10/10/96
31	lubber	8	55.5	58	103	11/12/96
58	rusty	10	35.0	22	101	10/10/96
58	rusty	10	35.0	58	103	11/12/96

$S \times R$



$\sigma_{S.sid < R.sid}(S \times R)$

Type of Join (Inner Join)

- **Theta join:** the general case of join is called Theta join
- **Equi join:** A special case of condition join where the condition c contains only *equalities*.

$$S \bowtie_{sid=sid} R$$

$$S \bowtie_{sid} R$$

$$S \bowtie R$$

S

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

R

<u>sid</u>	<u>bid</u>	<u>day</u>
22	101	10/10/96
58	103	11/12/96

$S \bowtie_{sid=sid} R$

(sid)	sname	rating	age	(sid)	bid	day
22	dustin	7	45.0	58	103	11/12/96
31	lubber	8	55.5	58	103	11/12/96

- **Natural join:** Equijoin on *all* common fields.

Example 1 of Algebra Query

- Find names of sailors who've reserved boat #103

<i>Sailors</i>	<u>sid</u>	sname	rating	age
	22	dustin	7	45.0
	31	lubber	8	55.5
	58	rusty	10	35.0

<i>Reserves</i>	<u>sid</u>	<u>bid</u>	<u>day</u>
	22	101	10/10/96
	58	103	11/12/96

❖ Answer

$$\pi_{sname}((\sigma_{bid=103} Reserves) \bowtie_{sid=sid} Sailors)$$

Internal
procedure:

$$\rho(Temp1, \sigma_{bid=103} Reserves)$$

$$\rho(Temp2, Temp1 \bowtie Sailors)$$

$$\pi_{sname}(Temp2) \quad * \text{ } \rho: \text{renaming operator}$$

❖ several alternative query expressions

$$\pi_{sname}((\sigma_{bid=103}(Reserves \bowtie_{sid=sid} Sailors))$$

Example 2

- Find names of sailors who've reserved a red boat

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

Sailors

<u>sid</u>	<u>bid</u>	day
22	101	10/10/96
22	102	10/10/98
58	103	11/12/96

Reserves

<u>bid</u>	bname	color
101	Interlake	blue
102	Interlake	red
103	Clipper	green

Boats

❖ Answer

$\pi_{sname}(((\sigma_{color='red'} Boats) \bowtie_{bid} Reserve) \bowtie_{sid} Sailors)$

❖ several alternative query expressions

$\pi_{sname}(\pi_{sid}((\pi_{bid} \sigma_{color='red'} Boats) \bowtie_{bid} Reserves) \bowtie_{sid} Sailors)$

Example 3

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

Sailors

<u>sid</u>	<u>bid</u>	day
22	101	10/10/96
22	102	10/10/98
58	103	11/12/96

Reserves

<u>bid</u>	bname	color
101	Interlake	blue
102	Interlake	red
103	Clipper	green

Boats

- Find sailors who've reserved a red or a green boat
 - Can identify all red or green boats, then find sailors who've reserved one of these boats:

❖ Answer

$$\pi_{sname}(((\sigma_{color='red' \text{ or } color='green'} Boats) \bowtie_{bid} Reserves) \bowtie_{sid} Sailors)$$

Internal procedure

$$\rho(Tempboats, (\sigma_{color='red' \vee color='green'} Boats))$$

$$\rho(Tempboats2, (Tempboats \bowtie Reserves))$$

$$\pi_{sname}(Tempboats2 \bowtie Sailors)$$

sname
dustin
rusty

Example 4

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

Sailors

<u>sid</u>	<u>bid</u>	day
22	101	10/10/96
22	102	10/10/98
58	103	11/12/96

Reserves

<u>bid</u>	bname	color
101	Interlake	blue
102	Interlake	red
103	Clipper	green

Boats

□ Find sailors who've reserved a red and a blue boat

- Previous approach won't work! Must identify sailors who've reserved red boats, sailors who've reserved blue boats, then find the intersection

~~$\pi_{sname}(((\sigma_{color='red' \text{ and } color='blue'} Boats) \bowtie_{bid} Reserves) \bowtie_{sid} Sailors)$~~

❖ Answer

$\pi_{sname}(((\pi_{sid}((\sigma_{color='red'} Boats) \bowtie_{bid} Reserves)) \cap (\pi_{sid}((\sigma_{color='blue'} Boats) \bowtie_{bid} Reserves))) \bowtie_{sid} Sailors)$

that is,

$\rho (Tempred, \pi_{sid}((\sigma_{color='red'} Boats) \bowtie Reserves))$

$\rho (Tempblue, \pi_{sid}((\sigma_{color='blue'} Boats) \bowtie Reserves))$

$\pi_{sname}((Tempred \cap Tempblue) \bowtie Sailors)$

sname
dustin